
SALMA: Structure-Aware Mask-Guided Alignment for Referring Image and Video Segmentation

Rui Peng^{1,2*} Xiongwei Lou^{1*} Ruitao Ma^{1,2}

Hao Liang¹ Xinyu Guo^{2,3†} Yinglun Li^{2,3†}

¹College of Mathematics and Computer Sciences, Zhejiang A & F University

²Beijing Research Center for Information Technology in Agriculture

³Beijing Key Laboratory of Digital Plant

Abstract

Unified multimodal large language models (MLLMs) have shown strong potential for referring image and video segmentation, but their cross-modal attention can drift toward salient distractors when queries require fine-grained spatial reasoning. We propose **SALMA**, a structure-aware alignment method that injects a class-agnostic mask prior into MLLM-based segmentation with minimal additional cost. SALMA performs a stop-gradient pre-pass through the SAM-2 mask decoder without positive user prompts, using SAM-2’s no-mask prompt embedding and dummy negative point convention to obtain low-resolution candidate-region logits. These logits are converted into a temperature-scaled spatial gate and used to modulate cross-attention outputs through a residual pathway with a learnable scale and warm-up schedule, preserving the pretrained model behavior during early optimization. We further introduce a text-mask contrastive objective and a boundary consistency loss to improve semantic and structural alignment. Experiments on image and video referring segmentation benchmarks show that SALMA improves over our reproduced Sa2VA-1B baseline, achieving 78.4 cIoU on RefCOCOg test and gains of +3.4 and +1.7 J&F on Ref-DAVIS17 and Ref-YouTube-VOS, respectively, with an approximately 0.7% latency increase. Code will be released upon acceptance.

1 Introduction

Referring segmentation—the task of localizing and delineating objects described by natural language—lies at the core of human-machine interaction. From embodied agents navigating cluttered environments (“pick up the red mug on the left”) to assistive systems for visually impaired users and precision agriculture [Luo et al., 2024], the ability to ground free-form language in pixel-precise masks is essential for actionable intelligence [Lai et al., 2024, Rasheed et al., 2024]. Early specialist models [Yang et al., 2022, Wang et al., 2022, Wu et al., 2022] achieved impressive results by designing task-specific architectures, yet they remain confined to task-specific benchmarks and lack the open-ended generalization demanded by real-world applications.

The recent surge of Multimodal Large Language Models (MLLMs) offers a promising alternative. By unifying visual perception with large-scale language reasoning, models like LLaVA [Liu et al., 2024a], InternVL [Chen et al., 2024], and Qwen-VL [Bai et al., 2023] demonstrate remarkable zero-shot capabilities across diverse vision-language tasks. Building on this foundation, unified grounding frameworks such as Sa2VA [Yuan et al., 2025] and LIRA [Li et al., 2025] integrate LLM reasoning with the pixel-level precision of SAM-2 [Ravi et al., 2024], aiming to enable general-purpose grounding within a single architecture.

*Equal contribution.

†Corresponding author.




Figure 1: **The Attention Hallucination Problem in Unified MLLMs.** When faced with complex spatial queries (e.g., depth, ordering, interaction), existing unified MLLMs such as OMG-LLaVA [Zhang et al., 2024b] (Center) exhibit severe attention drift, erroneously segmenting salient distractors rather than the queried target. This motivates our investigation into explicit spatial constraints. SALMA (Right) is our proposed solution.

However, a critical semantic-structural gap persists: existing MLLMs often treat visual features as abstract semantic tokens, weakening the low-level spatial priors needed for boundary delineation and instance-level disambiguation. This leads to *attention hallucination*—as shown in Figure 1, complex spatial queries (e.g., “the cup behind the laptop”) can cause attention to drift toward salient distractors rather than the spatially defined target. Our ablations in Table 2 indicate that explicitly biasing cross-modal attention with spatial priors is an effective way to reduce this failure mode in our setting.

To bridge this gap, we propose SALMA, a structure-aware method that re-introduces spatial fidelity into the attention mechanism. The key design choice is to reuse the SAM-2 decoder itself as a lightweight source of class-agnostic mask priors, rather than adding a separate localization head or relying on extra user prompts. By injecting these priors through Mask-Biased Attention inside the transformer layers, SALMA modulates cross-modal attention with dense structural cues, suppressing distractors while preserving the pretrained semantic representation.

Built upon this insight, our contributions can be summarized as follows:

- **Empirical Diagnosis: Spatial Constraint Deficiency.** We identify insufficient explicit spatial constraints in cross-modal attention as a plausible factor behind attention drift in unified MLLMs, supported by ablations and qualitative attention analysis.
- **Mask-Biased Attention (MBA).** We propose a soft residual attention-gating mechanism that repurposes the SAM-2 decoder through a stop-gradient pre-pass, using the segmentation decoder as a source of candidate-region priors without introducing an additional localization model.
- **Strong Performance with High Efficiency.** SALMA improves a Sa2VA-1B baseline across image and video referring segmentation benchmarks, achieving 71.9% J&F on Ref-DAVIS17 and 78.4 cIoU on RefCOCOg test [Pont-Tuset et al., 2017, Mao et al., 2016] with only a 0.7% latency increase.

2 Related Work

Large Multimodal Foundation Models. Large Language Models (LLMs) [OpenAI, 2023] have reshaped AI and motivated Multimodal LLMs that align visual encoders with language backbones. Early models such as LLaVA [Liu et al., 2024a] focus on static images, while later work scales resolution [Chen et al., 2024, Bai et al., 2023], strengthens reasoning [Li et al., 2024b], and extends to video [Lin et al., 2024, Bai et al., 2024]. Despite broad benchmarks [Liu et al., 2024b, Fu et al., 2023, Li et al., 2024a], many MLLMs still struggle with structure because images are processed as coarse semantic tokens.

Visual Segmentation and Grounding. Before unified MLLMs, referring segmentation was driven by specialist models, often with modular language–vision designs. Transformer-based methods place attention at the core of fusion; LAVT [Yang et al., 2022] and CRIS [Wang et al., 2022] inject language early and use contrastive alignment. For video, ReferFormer [Wu et al., 2022] treats RVOS as sequence prediction, while recent work strengthens grounding foundations [Liang et al., 2025], global-local video reasoning [Lin et al., 2025, Zheng et al., 2025], and interaction-aware expressions [Jin et al., 2025]. In parallel, universal segmentation advanced rapidly (e.g., SEEM [Zou et al., 2023], SegGPT [Wang et al., 2023]), and SAM-2 [Ravi et al., 2024] introduced promptable streaming memory. These models improve segmentation or temporal grounding, but they typically lack explicit LLM-level structural reasoning inside cross-modal attention.

Unified Multimodal Perception. To connect semantic reasoning with pixel grounding, recent work integrates segmentation decoders into MLLMs [Lai et al., 2024, Ren et al., 2024, Rasheed et al., 2024, Xia et al., 2024, Yuan et al., 2025, Li et al., 2025, Wei et al., 2025, Wang et al., 2024]. LISA [Lai et al., 2024] uses [SEG] tokens to invoke SAM-style decoding [Kirillov et al., 2023], PixelLM [Ren et al., 2024] adopts a lightweight pixel decoder and codebook, and GLaMM [Rasheed et al., 2024]/PSALM [Zhang et al., 2024c] add region- or mask-token interfaces. More recent systems broaden this line: InstructSeg [Wei et al., 2025] unifies instructed image and video segmentation, SegLLM [Wang et al., 2024] explores multi-round reasoning segmentation, Omni-RGPT [Heo et al., 2025] studies region-level image/video understanding via token marks, VideoGLaMM [Munasinghe et al., 2025] scales pixel-level video grounding, and UniPixel [Liu et al., 2025] connects object referring with pixel-level visual reasoning. These designs are powerful, but most pass segmentation intent through sparse tokens, prompts, or implicit alignment objectives. SALMA addresses a complementary problem by imposing a dense structural prior inside attention: SAM-2 supplies class-agnostic candidate regions, and the MLLM resolves which candidate satisfies the referring expression. Building on this analysis, we next detail the SALMA methodology.

3 Methodology

3.1 Problem Formulation

We build SALMA upon **Sa2VA** [Yuan et al., 2025], which integrates MLLM-based semantic reasoning with SAM-2 [Ravi et al., 2024] segmentation decoding, providing a strong baseline with open-ended generalization and foundational mask quality.

However, Sa2VA exhibits a structural bottleneck: while it excels at high-level semantic recognition, it degrades on fine-grained spatial instructions in two ways. (i) **Coarse localization:** the implicit [SEG] token injection produces fragmented masks when objects share similar attributes. (ii) **Structural hallucination:** although visual tokens retain positional information, cross-modal attention lacks an explicit dense structural prior, causing attention to drift toward distractors under complex spatial queries (e.g., relative positioning or heavy occlusion), as shown in Figure 1. Scaling data or parameters can improve overall accuracy, but it does not directly specify where semantic information should be injected into dense visual features. This motivates a complementary mechanism for re-introducing spatial priors, leading to our Mask-Biased Attention design.

3.2 Approach Overview

The overall pipeline of our proposed SALMA is illustrated in Figure 2. The approach consists of three integrated stages: (1) **Shared Visual Encoder:** A visual encoder and an LLM process the input video and text to extract multi-scale visual features and text embeddings. (2) **Mask-Biased Attention:** The




Figure 2: **Approach Overview.** (a) **Shared Visual Encoder.** Multi-scale visual features $\mathcal{F}_v$ are extracted by InternViT, while the LLM produces text embeddings $\mathcal{F}_t$ including the special [SEG] token h_{seg} . (b) **Mask-Biased Attention (MBA).** The *shared* SAM-2 decoder operates in a dual-mode fashion: a *stop-gradient pre-pass* (dashed arrows, stop-gradient) generates class-agnostic candidate-region logits $\mathcal{M}_{logits}$, which are transformed via a temperature-scaled Sigmoid ($\tau_{gate}=1.0$ in our final model) into a soft spatial gate G . This gate modulates ($\otimes$) the cross-attention output $\mathcal{O}_{attn}$ between visual queries and text keys/values, suppressing background noise while leaving target selection to the text-conditioned pathway. A deterministic linear warm-up coefficient $\omega(t)$ (rising from 0 to 1 over the first 5% of training steps) multiplies the learnable scale γ , so the effective gating strength is $\omega(t) \cdot \gamma$; this prevents the randomly initialized gate from disrupting pre-trained features in early training. (c) **Fine-grained Alignment.** The SAM-2 decoder, now with gradient tracking enabled, performs the final text-conditioned segmentation. Semantic alignment is enforced via TMC Loss ($\mathcal{L}_{tmc}$, contrastive matching between z_t and z_v^{mask}) and structural precision via Boundary Loss $\mathcal{L}_{bound}$ (Sobel edge matching). The total training objective combines $\mathcal{L}_{seg}$, $\mathcal{L}_{tmc}$, and $\mathcal{L}_{bound}$ with weights $\lambda_{tmc}=0.1$ and $\lambda_{bound}=0.05$.

proposed MBA mechanism injects pixel-level structural priors from the SAM-2 decoder into the cross-modal interaction to prevent attention drift. (3) **Fine-grained Alignment:** Finally, the SAM-2 decoder receives the text-conditioned features and predicts the target mask under a dual-constraint strategy involving TMC loss and Boundary Consistency loss.

3.3 Mask-Biased Attention

Standard cross-attention can introduce noise through unconstrained global interaction. We propose Mask-Biased Attention (MBA) to gate semantic injection using low-level spatial priors.

Spatial Prior via Auxiliary Execution. To obtain structural guidance without training a separate head, we leverage the *shared* SAM-2 decoder for a lightweight pre-pass. We invoke the decoder without positive user prompts: no box or mask is provided, and the implementation follows SAM-2’s standard no-mask embedding plus dummy negative-point convention. The resulting logits are used as soft candidate-region cues rather than target-specific masks; the referring expression is resolved by the subsequent text-conditioned attention and final decoder. This design is motivated by three considerations: (i) the SAM-2 decoder, pre-trained on SA-V, already encodes useful segmentation priors; (ii) the prompt-free pre-pass configuration is constant across training and inference, eliminating train-test distribution gaps; and (iii) the absence of user-supplied prompts keeps the prior purely data-driven. In our implementation, the pre-pass uses SAM-2’s language-free head output and detaches the resulting low-resolution logits from the computation graph before they are converted into the spatial gate. When multiple mask logits are produced, the implementation reduces them to a single gate-compatible logit map before MBA; in single-mask mode this is simply the sole output logit.

During final decoding, decoder weights are fine-tuned end-to-end. This dual-mode usage extracts useful priors with only a small runtime and memory overhead, quantified in Table 8.

Temporal Design for Video Inputs. For video referring segmentation, we adopt a *frame-independent prior with memory-consistent decoding* strategy. Specifically, the stop-gradient pre-pass generates spatial gates G independently for each frame, deliberately avoiding cross-frame propagation of the structural prior to prevent error accumulation from noisy early-frame predictions. Temporal coherence is instead handled by SAM-2’s built-in memory attention mechanism during the *final decoding* stage in Figure 2, where the decoder conditions on its memory bank to propagate identity-consistent masks across frames. This separation of concerns—frame-local structural gating for spatial precision, memory-based propagation for temporal consistency—allows MBA to complement rather than interfere with SAM-2’s native video capabilities.

Visual-Query Attention with Top-K Routing. Given the processed visual features $\mathcal{F}_v \in \mathbb{R}^{HW \times C}$ and text tokens $\mathcal{F}_t \in \mathbb{R}^{L \times C}$, we treat **visual features** $\mathcal{F}_v$ as **Queries** and **text tokens** $\mathcal{F}_t$ as **Keys/Values**, as illustrated in Figure 3. This design allows each spatial position on the dense visual grid to actively retrieve the most relevant semantic concepts from text, enabling fine-grained spatial-semantic alignment. To reduce computational redundancy, we employ a hard routing step that retains only a small subset of text tokens for cross-modal interaction. Specifically, the implementation ranks contextualized LLM hidden states with a lightweight two-layer scorer, $\text{Linear}(C, C/2) - \text{ReLU} - \text{Linear}(C/2, 1)$, and keeps the top $K=2$ tokens for cross-modal interaction. The scorer receives no separate supervision; it is used only to rank candidate text units before cross-attention. Since the top- K selection itself is discrete, we treat the selected indices as non-differentiable and optimize the selected text representations, attention projections, and downstream segmentation pathway through the standard training losses rather than through a continuous relaxation of the routing decision. We set $K=2$ as a practical default in this implementation. Let $\hat{\mathcal{F}}_t$ denote the gathered Top-K text tokens. For clarity, we present a simplified single-head form and omit per-head notation; the cross-modal attention map $A \in \mathbb{R}^{(HW) \times K}$ is computed as:

$$A = \text{Softmax} \left(\frac{(\mathcal{F}_v W_Q)(\hat{\mathcal{F}}_t W_K)^T}{\sqrt{d_k}} \right) \quad (1)$$

The retrieved semantic features are obtained by applying the attention weights to the filtered text values, e.g., $\mathcal{O}_{attn} = A(\hat{\mathcal{F}}_t W_V)W_O$ in the single-head notation above; our implementation uses the standard multi-head extension followed by the output projection.

Mask-Gated Modulation. We apply gating to the cross-attention output rather than the internal probability map, modulating semantic flow without distorting the probabilistic alignment between tokens, as illustrated in Figure 3.

We first normalize the logits $\mathcal{M}_{logits}$ into a soft spatial gate G using a temperature-scaled Sigmoid; with $\tau_{gate}=1.0$, this reduces to a standard Sigmoid in our final model. The structure-aware features are then injected into the visual backbone via a gated residual connection:

$$G = \sigma \left(\frac{\mathcal{M}_{logits}}{\tau_{gate}} \right), \quad \mathcal{F}'_v = \mathcal{F}_v + \gamma \cdot (\mathcal{O}_{attn} \odot G) \quad (2)$$

where G is resized via bilinear interpolation to match the spatial resolution $H \times W$ of the visual query features $\mathcal{F}_v$ and is broadcast along the channel dimension before multiplication. $\mathcal{O}_{attn}$ is the projected output of the Multi-Head Attention, $\odot$ denotes element-wise multiplication, and γ is a learnable scaling factor. To prevent the gate from perturbing pre-trained features before it has converged, we multiply γ by a deterministic warm-up coefficient $\omega(t)$ that linearly increases from 0 to 1 over the first 5% of training steps; the effective residual weight is therefore $\omega(t) \cdot \gamma$. After warm-up completes ($\omega=1$), the network freely adapts γ via gradient descent. This learnable γ controls the global strength of the residual structural prior. If the pre-pass mask is noisy or incorrect (e.g., in low-contrast camouflage), the spatial gate G limits where the residual can act, while optimization can reduce γ to fall back toward standard attention. Empirically, on RefCOCOg the effective residual magnitude, measured as the spatial mean of γG , has a median of 0.38 but drops below 0.05 when the pre-pass mask covers $<10\%$ of the target, suggesting that the residual branch can be down-weighted when the prior is weak.

Soft Constraint vs. Hard Tokens. A key distinction of our approach is the nature of the spatial injection. Existing methods like LISA Lai et al. [2024] or GLaMM Rasheed et al. [2024] rely




Figure 3: **The Mask-Biased Attention (MBA) Mechanism.** **Top:** Visual features $\mathcal{F}_v$ serve as Queries and text features as Keys/Values in cross-modal MHA. The resulting attention output $\mathcal{O}_{attn}$ is then element-wise multiplied ($\odot$) with a soft spatial gate $G = \sigma(\mathcal{M}/\tau)$. Crucially, this gate is derived from the *stop-gradient pre-pass* of the SAM-2 decoder, which provides class-agnostic structural priors without additional training. The gated features are injected back via a learnable residual: $\mathcal{F}'_v = \mathcal{F}_v + \gamma \cdot (\mathcal{O}_{attn} \odot G)$. **Bottom:** The resulting *spotlight effect* biases semantic injection toward object-like regions proposed by the pre-pass, reducing background activation while preserving a residual path for text-conditioned target selection.

on “hard” tokens (e.g., [SEG]) to trigger segmentation, creating a representational bottleneck: the complex referential intent must be compressed into a single vector, often leading to “all-or-nothing” failures if the token embedding is misaligned. In contrast, MBA imposes a *soft, dense constraint* over the entire visual feature map. The residual connection $\mathcal{F}'_v = \mathcal{F}_v + \dots$ helps preserve the original semantic information rather than discarding it. If the prior is incorrect, the learnable γ can reduce the influence of the residual branch, allowing the model to remain closer to standard attention. Moreover, dense gating preserves fine-grained geometric cues for small or thin objects that may be weakened in token-based bottlenecks.

3.4 Fine-grained Alignment

MBA injects spatial priors, but precise masks still benefit from explicit fine-grained supervision. Without targeted constraints, features can drift to semantically similar yet spatially distinct instances (e.g., adjacent “red cars”). We therefore apply a dual-constraint strategy (TMC + boundary), as shown in the **Right** panel of Figure 2.

Text-Mask Contrastive Learning. To prevent the visual features from drifting away from the textual instructions during the decoding process, we enforce a semantic consistency constraint inspired by contrastive vision-language alignment Radford et al. [2021]. Specifically, we extract the region-level visual feature z_v^{mask} by average-pooling the post-MBA feature map $\mathcal{F}'_v$ over the ground-truth foreground region $\mathcal{M}^{gt}$. We use the ground-truth mask during training to promote stable semantic alignment, decoupling it from early-stage segmentation errors. Let z_t be the embedding of the special [SEG] token corresponding to the referring expression. We employ a Symmetric InfoNCE van den Oord et al. [2018] loss where *negatives are constructed from other text-image pairs within the same batch*, maximizing the mutual information between the matched text-visual pair (z_t, z_v^{mask}) :

$$\mathcal{L}_{tmc} = \mathcal{L}_{v \rightarrow t} + \mathcal{L}_{t \rightarrow v} \quad (3)$$

where both directions (vision-to-text and text-to-vision) are normalized by a fixed temperature $\tau_{tmc}=0.15$.

Table 1: **Comparison with Representative Unified MLLMs and Segmentation Methods.** We report cIoU for image referring segmentation tasks, $\mathcal{J}\&\mathcal{F}$ for video segmentation tasks, and standard metrics for multimodal understanding. Results for prior methods are from their original papers unless marked by *, so this table provides a contextual comparison across representative systems rather than a fully controlled re-training study. Bold numbers highlight the best reported results within each metric group when the comparison is meaningful under the shared table format; the reproduced Sa2VA-1B row is the primary controlled baseline for SALMA. * denotes our reproduction under the same training data and hardware setup for fair comparison with SALMA. A dash (–) indicates unsupported tasks or unreported results.

Method	Image Segmentation						Video Segmentation				Image Chat			
	RefCOCO			RefCOCO+			RefCOCOg		Ref-DAVIS17	Ref-YT-VOS	MeVis	MME	MMB	SEED
	Val	TestA	TestB	Val	TestA	TestB	Val	Test						
LLaVA-1.5-13B Liu et al. [2024a]	-	-	-	-	-	-	-	-	-	-	-	1531(+)	68.8	70.1
Video-LLaVA-7B Lin et al. [2024]	-	-	-	-	-	-	-	-	-	-	-	-	60.9	-
LLaMA-VID-7B Li et al. [2023]	-	-	-	-	-	-	-	-	-	-	-	1521(+)	65.1	59.9
mPLUG-Owl3-8B Ye et al. [2024]	-	-	-	-	-	-	-	-	-	-	-	-	77.6	-
InternVL2-8B Chen et al. [2024]	-	-	-	-	-	-	-	-	-	-	-	-	81.7	76.2
PixelLM-7B Ren et al. [2024]	73.0	-	-	66.3	-	-	69.3	-	-	-	-	309/135	17.4	-
PixelLM-13B Ren et al. [2024]	73.0	76.5	68.2	66.3	71.7	58.3	69.3	70.5	-	-	-	309/135	17.4	-
LaSagnA Wei et al. [2024]	76.8	-	-	66.4	-	-	70.6	-	-	-	-	0/0	0.0	-
LISA-7B Lai et al. [2024]	74.1	76.5	71.1	62.4	67.4	56.5	66.4	68.5	-	-	-	1/1	0.4	-
LISA-GLEE Wu et al. [2024]	76.4	78.2	73.8	67.3	71.3	62.3	71.6	72.4	-	-	-	-	-	-
GLaMM Rasheed et al. [2024]	79.5	83.2	76.9	72.6	78.7	64.6	74.2	74.9	-	-	-	14/9	36.8	-
LLaVA-G-7B Zhang et al. [2024a]	77.1	-	-	68.8	-	-	71.5	-	-	-	-	-	-	-
GSVA Xia et al. [2024]	77.2	78.9	73.5	65.9	69.6	59.8	72.7	73.3	-	-	-	-	-	-
SegLLM Wang et al. [2024]	80.2	81.5	75.4	70.3	73.0	62.5	72.6	73.6	-	-	-	-	-	-
OMG-LLaVA-8B Zhang et al. [2024b]	75.6	77.7	71.2	65.6	69.7	58.9	70.7	70.2	-	-	-	1177/235	47.9	56.5
PSALM Zhang et al. [2024c]	83.6	84.7	81.6	72.9	75.5	70.1	73.8	74.4	-	-	-	-	52.5	-
VideoLISA-3.8B Bai et al. [2024]	73.8	-	-	63.4	-	-	68.3	-	68.8	63.7	44.4	-	-	-
VISA-13B Yan et al. [2024]	72.4	-	-	59.8	-	-	65.5	-	70.4	63.0	44.5	-	-	-
LIRA-2B Li et al. [2025]	79.5	81.9	76.0	72.6	77.4	66.1	75.4	75.1	-	-	-	-	74.0	-
ReferDINO Liang et al. [2025]	-	-	-	-	-	-	-	-	68.9	69.3	-	-	-	-
Sa2VA-1B (Base)*	79.6	82.7	76.5	73.6	79.0	66.6	77.8	77.5	68.5	65.3	41.7	1487/433	71.9	71.0
SALMA	80.4	83.2	77.7	74.8	80.0	68.9	78.0	78.4	71.9	67.0	46.0	1421/353	71.2	69.9

Boundary Consistency Constraint. Standard binary cross-entropy (BCE) and Dice losses primarily focus on the overall mask area but are insensitive to boundary errors. In video segmentation, however, fuzzy boundaries are a major source of qualitative degradation (e.g., temporal jitter). We incorporate a Boundary Loss $\mathcal{L}_{bound}$ that measures the discrepancy between predicted and ground-truth edge maps:

$$\mathcal{L}_{bound} = \|\widetilde{\nabla \hat{\mathcal{M}}} - \widetilde{\nabla \mathcal{M}^{gt}}\|_1 \quad (4)$$

where $\nabla \hat{\mathcal{M}}$ and $\nabla \mathcal{M}^{gt}$ denote the gradient magnitudes of the predicted and ground-truth masks, respectively, computed via Sobel edge detection with 3×3 kernels, and $\widetilde{\cdot}$ denotes per-sample min-max normalization applied to the edge maps for numerical stability. During training, $\mathcal{L}_{bound}$ is applied independently to every sampled frame (both image and video), ensuring that per-frame boundary precision is directly supervised rather than relying solely on temporal propagation. By optimizing this boundary-specific objective, the model is explicitly penalized for edge inaccuracies.

3.5 Unified Training Objective

We formulate a unified training objective that balances global semantic fidelity with local structural precision:

$$\mathcal{L}_{total}(t) = \mathcal{L}_{seg} + \omega(t)\lambda_{tmc}\mathcal{L}_{tmc} + \omega(t)\lambda_{bound}\mathcal{L}_{bound} \quad (5)$$

where:

- $\mathcal{L}_{seg} = \mathcal{L}_{dice} + \mathcal{L}_{bce}$: This is the standard linear combination of Dice loss and Binary Cross-Entropy loss used in SAM-2, responsible for the general mask quality.
- $\mathcal{L}_{tmc}$: The Text-Mask Contrastive loss enforces semantic fidelity.
- $\mathcal{L}_{bound}$: The Boundary loss enforces high-frequency structural precision.

4 Experiments

4.1 Implementation Details

Architecture. We build on Sa2VA-1B, using InternVL2.5-1B Chen et al. [2025] (with InternViT encoder) as the MLLM backbone and SAM-2’s Hiera-L as the segmentation image encoder. The SAM-2 mask decoder serves dual roles: generating detached candidate-region logits during a stop-gradient pre-pass for prior extraction, and performing full-gradient text-conditioned decoding.

Training. We train end-to-end for 1 epoch on $4\times$ RTX 5090 GPUs (per-GPU batch size 2, gradient accumulation 4) using AdamW ($\beta_1=0.9$, $\beta_2=0.999$, weight decay 0.05), learning rate 4×10^{-5} with cosine decay, DeepSpeed ZeRO-2, and BF16 mixed precision. The training mixture covers four data families: image referring segmentation (RefCOCO+/g Kazemzadeh et al. [2014], Mao et al. [2016]), video referring segmentation (MeVis, ReVOS Yan et al. [2024], Ref-YouTube-VOS/Ref-DAVIS Xu et al. [2018], Pont-Tuset et al. [2017]), grounded conversation/region understanding (GCG Rasheed et al. [2024], Osprey-724k Yuan et al. [2024]), and general image/video instruction tuning (LLaVA-1.5 Liu et al. [2024a], Video-ChatUniVi Jin et al. [2024]). To keep structurally demanding segmentation examples visible during mixed training, we up-sample RefCOCO+/g, MeVis, Ref-YouTube-VOS, Ref-DAVIS, and SAM-2 SA-V by $4\times$, ReVOS by $10\times$, and keep general instruction data at $1\times$; the full dataset table is provided in Table 3. For the class-agnostic SA-V subset, we use our Ref-SAV formatting: each selected masklet is paired with a long object description or sampled short caption, inserted into a segmentation prompt, and supervised with the mask sequence and a [SEG] response. Official train splits are used for all video benchmarks.

Hyper-parameters. We set $\lambda_{tmc}=0.1$, $\lambda_{bound}=0.05$, and $\tau_{gate}=1.0$. A 5% linear warm-up (coefficient ω , see Section 3.3) is applied to auxiliary loss weights and the gating scale γ to prevent disrupting pre-trained features.

4.2 Evaluation Benchmarks

We evaluate across three complementary domains:

Referring Image Segmentation. We evaluate on the classic RefCOCO, RefCOCO+, and RefCOCOg benchmarks. These datasets test the model’s ability to ground static objects from natural language descriptions. RefCOCOg is particularly challenging due to its longer, more complex expressions. We report performance using the standard Cumulative Intersection-over-Union metric.

Referring Video Segmentation. To assess structural stability in dynamic scenes, we test on Ref-DAVIS17 and Ref-YouTube-VOS. These benchmarks require the model to maintain consistent tracking of a referred object across video frames. We report the $\mathcal{J}\&\mathcal{F}$ mean score, which aggregates region similarity ($\mathcal{J}$) and boundary contour accuracy ($\mathcal{F}$).

Motion-Centric Video Segmentation. We utilize the MeVis benchmark, which features queries explicitly dependent on motion cues (e.g., “the fish swimming to the left”). This tests whether our structural priors remain robust under complex temporal dynamics.

4.3 Comparison with Leading MLLMs

Quantitative Analysis. As shown in Table 1, SALMA improves over the reproduced Sa2VA-1B baseline on all evaluated referring segmentation benchmarks and remains competitive on general multimodal understanding metrics. Since most prior results are copied from original papers with different model scales, training data, and evaluation details, we use them as contextual references and use the reproduced Sa2VA-1B row as the primary controlled comparison. Gains over this baseline are largest in structurally demanding settings: +3.4 J&F on Ref-DAVIS17, +4.3 on MeVis, and +1.7 on Ref-YouTube-VOS, with 78.0/78.4 cloU on RefCOCOg val/test. Recent methods highlight the breadth-specialization trade-off: SegLLM Wang et al. [2024] improves single-round image referring segmentation while primarily targeting multi-round interaction, and ReferDINO Liang et al. [2025] achieves the highest Ref-YouTube-VOS score as a dedicated RVOS method. Larger or more specialized recent systems such as InstructSeg Wei et al. [2025] are compared separately in Table 5 because their architectures and training recipes are not directly controlled against our 1B setting. SALMA remains competitive on video grounding, reports strong Ref-DAVIS17 and MeVis scores among the listed unified MLLM-style models, and preserves general multimodal performance

Table 2: **Ablation Study of Components.** We progressively add Mask-Biased Attention, Text-Mask Contrastive and Boundary Loss.

Configuration	DAVIS (J&F)	RefCOCOg (test)
Baseline (Sa2VA)	68.52	77.45
+ MBA	69.83	78.21
+ MBA + TMC	71.04	78.28
+ MBA + TMC + Bound	71.87	78.42

without relying on a video-specialist architecture. The stronger video gains suggest that accurate per-frame structural priors complement SAM-2’s temporal memory, reducing the likelihood that early attention drift propagates across frames.

Generalization-Specialization Trade-off. SALMA shows small drops on general MLLM benchmarks (MME $-66/-80$, MMB -0.7 , SEED -1.1), suggesting a mild shift of capacity toward grounding under the same training mixture and hardware setup. The decrease remains limited compared with hard-token segmentation systems such as PSALM Zhang et al. [2024c] (52.5 MMBench vs. 71.2 for SALMA), indicating that soft residual gating better preserves general reasoning behavior.

4.4 Ablation Study

We conduct a progressive ablation in Table 2. MBA provides the largest first-step gain (+1.31 J&F on DAVIS and +0.76 cIoU on RefCOCOg), while TMC and Boundary loss add complementary improvements, especially for video segmentation and boundary precision. The full model performs best (71.87 J&F / 78.42 cIoU), supporting the dual-constraint design.

Efficiency. On 50 Ref-DAVIS17 validation clips with a single RTX 5090 in BF16, SALMA runs at 17.84 FPS versus 17.97 FPS for Sa2VA-1B ($\sim 0.7\%$ latency increase). The stop-gradient pre-pass adds roughly 0.4 ms per frame, <1 GFLOP, and +13 MB peak GPU memory.

Scaling and Additional Analysis. SALMA also improves Sa2VA-4B on all RefCOCO+/g splits, with the largest gain on RefCOCOg Val (+0.9 cIoU). Appendix results further cover FiLM ablations, saliency bias, computational cost, and qualitative examples.

Limitations. SALMA depends on the class-agnostic SAM-2 prior and is less effective on extreme micro-objects ($< 1\%$ area), dense near-identical instances, and abstract physical-state reasoning. It also introduces a small pre-pass overhead; future work may combine the gate with active zoom-in or attribute-aware routing.

5 Conclusion

We identify insufficient spatial constraints as a bottleneck in unified MLLMs for referring segmentation and introduce SALMA, which repurposes the SAM-2 decoder as a class-agnostic prior generator. Through stop-gradient pre-pass gating, text-mask contrastive alignment, and boundary supervision, SALMA injects dense structural bias while preserving the pretrained segmentation pathway. With only $\sim 0.7\%$ latency overhead, SALMA improves over Sa2VA-1B across image and video benchmarks, with particularly strong gains on Ref-DAVIS17, MeVis, and Ref-YouTube-VOS. Ablations show that explicit spatial gating is the largest contributor. These results suggest that lightweight structural priors can improve foundation-model grounding without replacing pretrained semantic reasoning.

References

- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-VL: A versatile vision-language model for understanding, localization, text reading, and beyond. *arXiv preprint arXiv:2308.12966*, 2023.
- Zechen Bai, Tong He, Haiyang Mei, Pichao Wang, Ziteng Gao, Joya Chen, Lei Liu, Zheng Zhang, and Mike Z Shou. One token to seg them all: Language instructed reasoning segmentation in videos. *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

- Zhe Chen, Jiannan Wu, Wenhai Wang, Weijie Su, Guo Chen, Sen Xing, Muyan Zhong, Qinglong Zhang, Xizhou Zhu, Lewei Lu, Bin Li, Ping Luo, Tong Lu, Yu Qiao, and Jifeng Dai. InternVL: Scaling up vision foundation models and aligning for generic visual-linguistic tasks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Zhe Chen, Weiyun Wang, Yue Cao, Yangzhou Liu, Zhangwei Gao, Erfei Cui, Jinguo Zhu, Shenglong Ye, Hao Tian, Zhaoyang Liu, Lixin Gu, Xuehui Wang, Qingyun Li, Yiming Ren, Zixuan Chen, Jiapeng Luo, Jiahao Wang, Tan Jiang, Bo Wang, Conghui He, Botian Shi, Xingcheng Zhang, Han Lv, Yi Wang, Wenqi Shao, Pei Chu, Zhongying Tu, Tong He, Zhiyong Wu, Huipeng Deng, Jiaye Ge, Kai Chen, Kaipeng Zhang, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhai Wang. Expanding performance boundaries of open-source multimodal models with model, data, and test-time scaling. *arXiv preprint arXiv:2412.05271*, 2025.
- Chaoyou Fu, Peixian Chen, Yunhang Shen, Yulei Qin, Mengdan Zhang, Xu Lin, Jinrui Yang, Xiawu Zheng, Ke Li, Xing Sun, et al. MME: A comprehensive evaluation benchmark for multimodal large language models. *arXiv preprint arXiv:2306.13394*, 2023.
- Miran Heo, Min-Hung Chen, De-An Huang, Sifei Liu, Subhashree Radhakrishnan, Seon Joo Kim, Yu-Chiang Frank Wang, and Ryo Hachiuma. Omni-RGPT: Unifying image and video region-level understanding via token marks. *arXiv preprint arXiv:2501.08326*, 2025.
- Peng Jin, Ryuichi Takanobu, Wancai Zhang, Xiaochun Cao, and Li Yuan. Chat-UniVi: Unified visual representation empowers large language models with image and video understanding. *arXiv preprint arXiv:2311.08046*, 2024.
- Woojeong Jin, Seongchan Kim, Jaeho Lee, and Seungryong Kim. InterRVOS: Interaction-aware referring video object segmentation. *arXiv preprint arXiv:2506.02356*, 2025.
- Sahar Kazemzadeh, Vicente Ordonez, Mark Matten, and Tamara Berg. ReferItGame: Referring to objects in photographs of natural scenes. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2014.
- Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C Berg, Wan-Yen Lo, et al. Segment anything. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- Xin Lai, Zhuotao Tian, Yukang Chen, Yanwei Li, Yuhui Yuan, Shu Liu, and Jiaya Jia. LISA: Reasoning segmentation via large language model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Bohao Li, Yuying Ge, Yixiao Ge, Guangzhi Wang, Rui Wang, Ruimao Zhang, and Ying Shan. SEED-Bench: Benchmarking multimodal large language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024a.
- Feng Li, Renrui Zhang, Hao Zhang, Yuanhan Zhang, Bo Li, Wei Li, Zejun Ma, and Chunyuan Li. LLaVA-NeXT-Interleave: Tackling multi-image, video, and 3d in large multimodal models. *arXiv preprint arXiv:2407.07895*, 2024b.
- Yanwei Li, Chengyao Wang, and Jiaya Jia. LLaMA-VID: An image is worth 2 tokens in large language models. *arXiv preprint arXiv:2311.17043*, 2023.
- Zhang Li, Biao Yang, Qiang Liu, Shuo Zhang, Zhiyin Ma, Liang Yin, Linger Deng, Yabo Sun, Yuliang Liu, and Xiang Bai. LIRA: Inferring segmentation in large multi-modal models with local interleaved region assistance. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025.
- Tianming Liang, Kun-Yu Lin, Chaolei Tan, Jianguo Zhang, Wei-Shi Zheng, and Jian-Fang Hu. ReferDINO: Referring video object segmentation with visual grounding foundations. *arXiv preprint arXiv:2501.14607*, 2025.
- Bin Lin, Yang Ye, Bin Zhu, Jiayi Cui, Munan Ning, Peng Jin, and Li Yuan. Video-LLaVA: Learning united visual representation by alignment before projection. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 5971–5984, 2024.

- Lang Lin, Xueyang Yu, Ziqi Pang, and Yu-Xiong Wang. GLUS: Global-local reasoning unified into a single large language model for video segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 8658–8667, 2025.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024a.
- Ye Liu, Zongyang Ma, Junfu Pu, Zhongang Qi, Yang Wu, Ying Shan, and Chang Wen Chen. UniPixel: Unified object referring and segmentation for pixel-level visual reasoning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=3901vVCfo3>.
- Yuan Liu, Haodong Duan, Yuanhan Zhang, Bo Li, Songyang Zhang, Wangbo Zhao, Yike Yuan, Jiaqi Wang, Conghui He, Ziwei Liu, et al. MMBench: Is your multi-modal model an all-around player? In *European Conference on Computer Vision (ECCV)*, 2024b.
- Zifei Luo, Wenzhu Yang, Yunfeng Yuan, Ruru Gou, and Xiaonan Li. Semantic segmentation of agricultural images: A survey. *Information Processing in Agriculture*, 2024.
- Junhua Mao, Jonathan Huang, Alexander Toshev, Oana Camburu, Alan L Yuille, and Kevin Murphy. Generation and comprehension of unambiguous object descriptions. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- Shehan Munasinghe, Hanan Gani, Wenqi Zhu, Jiale Cao, Eric Xing, Fahad Shahbaz Khan, and Salman Khan. VideoGLaMM: A large multimodal model for pixel-level visual grounding in videos. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 19036–19046, 2025.
- OpenAI. GPT-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Ethan Perez, Florian Strub, Harm De Vries, Vincent Dumoulin, and Aaron Courville. FiLM: Visual reasoning with a general conditioning layer. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2018.
- Jordi Pont-Tuset, Federico Perazzi, Sergi Caelles, Pablo Arbeláez, Alex Sorkine-Hornung, and Luc Van Gool. The 2017 davis challenge on video object segmentation. *arXiv preprint arXiv:1704.00675*, 2017.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2021.
- Hanoona Rasheed, Muhammad Maaz, Sahal Shaji, Abdelrahman Shaker, Salman Khan, Hisham Cholakkal, Rao M Anwer, Eric Xing, Ming-Hsuan Yang, and Fahad S Khan. GLaMM: Pixel grounding large multimodal model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Nikhila Ravi, Valentin Gabeur, Yuan-Ting Hu, Ronghang Hu, Chaitanya Ryali, Tengyu Ma, Haitham Khedr, Roman Rädle, Chloe Rolland, Laura Gustafson, et al. SAM 2: Segment anything in images and videos. *arXiv preprint arXiv:2408.00714*, 2024.
- Zhongwei Ren, Zhicheng Huang, Yunchao Wei, Yao Zhao, Dongmei Fu, Jiashi Feng, and Xiaoje Jin. PixelLM: Pixel reasoning with large multimodal model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- Xinlong Wang, Xiaosong Zhang, Yue Cao, Wen Wang, Chunhua Shen, and Tiejun Huang. SegGPT: Segmenting everything in context. *arXiv preprint arXiv:2304.03284*, 2023.

- XuDong Wang, Shaolun Zhang, Shufan Li, Konstantinos Kallidromitis, Kehan Li, Yusuke Kato, Kazuki Kozuka, and Trevor Darrell. SegLLM: Multi-round reasoning segmentation. *arXiv preprint arXiv:2410.18923*, 2024.
- Zhaoqing Wang, Yu Lu, Qiang Li, Xunqiang Tao, Yandong Guo, Mingming Gong, and Tongliang Liu. CRIS: CLIP-driven referring image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Cong Wei, Haoxian Tan, Yujie Zhong, Yujiu Yang, and Lin Ma. LaSagna: Language-based segmentation assistant for complex queries. *arXiv preprint arXiv:2404.08506*, 2024.
- Cong Wei, Yujie Zhong, Haoxian Tan, Yingsen Zeng, Yong Liu, Hongfa Wang, and Yujiu Yang. InstructSeg: Unifying instructed visual segmentation with multi-modal large language models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 20193–20203, 2025.
- Jiannan Wu, Yi Jiang, Peize Sun, Zehuan Yuan, and Ping Luo. Language as queries for referring video object segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Junfeng Wu, Yi Jiang, Qihao Liu, Zehuan Yuan, Xiang Bai, and Song Bai. General object foundation model for images and videos at scale. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Zhuofan Xia, Dongchen Han, Yizeng Han, Xuran Pan, Shiji Song, and Gao Huang. GSVA: Generalized segmentation via multimodal large language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Ning Xu, Linjie Yang, Yuchen Fan, Dingcheng Yue, Yuchen Liang, Jianchao Yang, and Thomas Huang. YouTube-VOS: A large-scale video object segmentation benchmark. *arXiv preprint arXiv:1809.03327*, 2018.
- Cilin Yan, Haochen Wang, Shilin Yan, Xiaolong Jiang, Yao Hu, Guoliang Kang, Weidi Xie, and Efstratios Gavves. ViSA: Reasoning video object segmentation via large language models. In *European Conference on Computer Vision (ECCV)*, 2024.
- Zhao Yang, Jiaqi Wang, Yansong Tang, Kai Chen, Hengshuang Zhao, and Philip HS Torr. LAVT: Language-aware vision transformer for referring image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Jiabo Ye, Haiyang Xu, Haowei Liu, Anwen Hu, Ming Yan, Qi Qian, Ji Zhang, Fei Huang, and Jingren Zhou. mPLUG-Owl3: Towards long image-sequence understanding in multi-modal large language models. *arXiv preprint arXiv:2408.04840*, 2024.
- Haobo Yuan, Xiangtai Li, Tao Zhang, Yueyi Sun, Zilong Huang, Shilin Xu, Shunping Ji, Yunhai Tong, Lu Qi, Jiashi Feng, and Ming-Hsuan Yang. Sa2VA: Marrying SAM2 with LLaVA for dense grounded understanding of images and videos. *arXiv preprint arXiv:2501.04001*, 2025.
- Yuqian Yuan, Wentong Li, Jian Liu, Dongqi Tang, Xinjie Luo, Chi Qin, Lei Zhang, and Jianke Zhu. Osprey: Pixel understanding with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Hao Zhang, Hongyang Li, Feng Li, Tianhe Ren, Xueyan Zou, Shilong Liu, Shijia Huang, Jianfeng Gao, Lei Zhang, Chunyuan Li, et al. LLaVA-Grounding: Grounded visual chat with large multimodal models. *European Conference on Computer Vision (ECCV)*, 2024a.
- Tao Zhang, Xiangtai Li, Hao Fei, Haobo Yuan, Shengqiong Wu, Shunping Ji, Chen Change Loy, and Shuicheng Yan. OMG-LLaVA: Bridging image-level, object-level, pixel-level reasoning and understanding. *Advances in Neural Information Processing Systems (NeurIPS)*, 2024b.
- Zheng Zhang, Yeyao Ma, Enming Zhang, and Xiang Bai. PSALM: Pixelwise segmentation with large multi-modal model. *European Conference on Computer Vision (ECCV)*, 2024c.

- Rongkun Zheng, Lu Qi, Xi Chen, Yi Wang, Kun Wang, and Hengshuang Zhao. ViLLa: Video reasoning segmentation with large language model. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 23667–23677, 2025.
- Xueyan Zou, Jianwei Yang, Hao Zhang, Feng Li, Linjie Li, Jianfeng Wang, Lijuan Wang, Jianfeng Gao, and Yong Jae Lee. Segment everything everywhere all at once. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

A Additional Analyses and Qualitative Results

This appendix contains supporting material for the compact summaries in the main text, including the training mixture, model scaling results, additional ablations, computational analysis, and qualitative examples.

A.1 Training Mixture Details

Table 3: **Training Data Mixture.** Datasets used during SALMA training. Repeat factors are used to keep structurally demanding segmentation samples visible during mixed training.

Dataset	Task Type	Approx. Samples	Repeats	Primary Purpose
RefCOCO / RefCOCO+	Image RIS	~100k	4×	Foundational image grounding
RefCOCOg	Image RIS	~80k	4×	Complex spatial/linguistic reasoning
RefCOCOg-GCG	Image GCG	~80k	1×	Grounded conversation generation
GranDf	Image GCG	~30k	10×	Fine-grained grounded captioning
MeVis	Video RVOS	~6k videos	4×	Motion-centric video segmentation
ReVOS	Video RVOS	~5k videos	10×	Reasoning video object segmentation
Ref-YouTube-VOS	Video RVOS	~3k videos	4×	Standard video RIS
SAM-2 SA-V (Ref-SAV subset)	Pseudo Video RIS	~50k videos	4×	Structural prior and mask robustness
Osprey-724k	Region Conversation	~724k	1×	Region-level multimodal instruction
LLaVA-1.5 / Video-ChatUniVi	General VQA/Chat	Mixed	1×	General MLLM ability retention

A.2 Model Scaling

Table 4: **Scaling to 4B Parameters.** cIoU on RefCOCO, RefCOCO+, and RefCOCOg. SALMA-4B consistently outperforms Sa2VA-4B across all splits.

Method	RefCOCO			RefCOCO+			RefCOCOg	
	Val	TestA	TestB	Val	TestA	TestB	Val	Test
Sa2VA-4B	82.4	84.3	79.5	77.6	81.2	73.1	79.7	80.4
SALMA-4B	82.6	84.9	80.0	78.1	82.0	73.4	80.6	81.1
Δ	+0.2	+0.6	+0.5	+0.5	+0.8	+0.3	+0.9	+0.7

Table 5: **Larger-System Comparison.** InstructSeg is reported separately because it uses a different multi-component instructed-segmentation architecture and training recipe, while the main comparison focuses on the 1B SALMA model and representative comparable baselines. A dash indicates that the metric is not reported under the shared table format.

Method	Scale	RefCOCO			RefCOCO+			RefCOCOg		Video Seg.			MMB
		Val	TestA	TestB	Val	TestA	TestB	Val	Test	Ref-DAVIS17	Ref-YT-VOS	MeVis	
SALMA	1B	80.4	83.2	77.7	74.8	80.0	68.9	78.0	78.4	71.9	67.0	46.0	71.2
SALMA	4B	82.6	84.9	80.0	78.1	82.0	73.4	80.6	81.1	-	-	-	-
InstructSeg	2.7B+	85.8	86.6	84.0	80.1	83.8	75.6	79.3	80.3	71.1	67.5	-	68.6

The 4B results show that the SALMA gain persists as model capacity increases, with positive improvements on every image referring segmentation split. InstructSeg obtains stronger image RIS scores with a different instructed-segmentation system, while SALMA-1B targets a lower-cost operating point and remains competitive on video grounding.

A.3 Feature Modulation Ablation

We compared against Residual FiLM Perez et al. [2018], the canonical feature modulation paradigm ($\gamma \cdot \mathcal{F} + \beta$). Replacing our gating with FiLM degrades performance on this ablation, suggesting that full affine modulation may distort the pre-trained feature statistics used by SAM-2.

Table 6: **Analysis of Feature Modulation.** Comparing our simplified spatial gating vs. complex FiLM modulation. MeVis is evaluated on the Val_u split, which differs from the MeVis test-set result reported in Table 1. Values in parentheses indicate a performance drop relative to our full model.

Method	Ref-DAVIS17	MeVis (Val_u)
Ours (Full)	71.87	53.37
Ours (w/ FiLM)	68.13 (-3.74)	47.46 (-5.91)

Table 7: **Saliency Bias Analysis on RefCOCOg (Val).** We report cIoU scores grouped cumulatively by object area ratio.

Method	$< 1\%$ ($N = 35$)	$< 2\%$ ($N = 151$)	$< 5\%$ ($N = 1559$)	All ($N = 4896$)
Baseline (Sa2VA)	37.1	47.7	68.0	77.8
Ours	33.2	48.9	70.2	78.0

A.4 Saliency Bias Analysis

SALMA improves on objects occupying $< 2\%$ and $< 5\%$ of the image area, while dropping on the extreme $< 1\%$ subset. This subset accounts for only 0.7% of the validation samples, but it identifies a useful failure mode for future high-resolution prompting.

A.5 Computational Cost

Table 8: **Computational Cost Analysis on 1024×1024 resolution.** The SALMA pre-pass introduces negligible overhead relative to the backbone encoders.

Component	Params (M)	FLOPs (G)
Vision Encoder (InternViT)	304.0	1614.1
SAM-2 Image Encoder (Hiera)	212.7	810.5
Visual Projector	4.5	1.2
SALMA Pre-Pass (Gating + Decoder)	12.0	< 1.0

A.6 Qualitative Results




Figure 4: **Qualitative Comparison on Complex Referring Expressions.** We present six diverse examples comparing the Baseline (Sa2VA-1B) with SALMA.




Figure 5: **Spatial gate maps (ours) vs. cross-modal attention (baseline).** Mask-Biased Attention helps mitigate attention drift under complex referring expressions.




Figure 6: **Qualitative Results across Diverse Scenarios.** SALMA improves temporal consistency, small-object grounding, and complex spatial reasoning.